

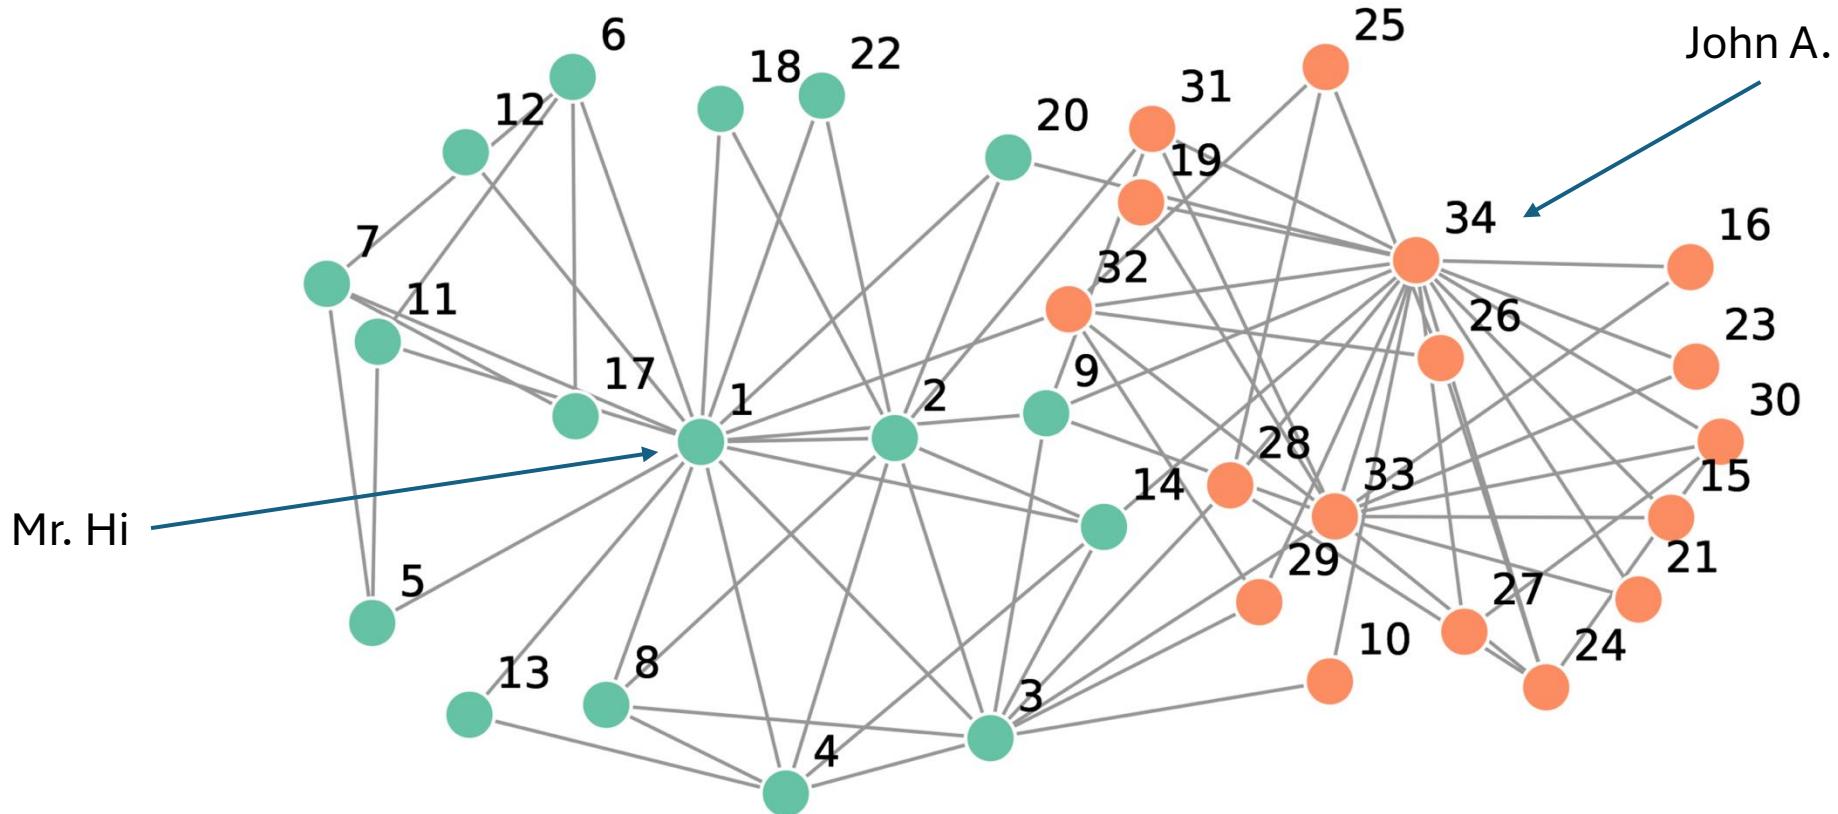
Graph Based Learning

Clustering and Dimensionality Reduction
Methods in Data Science
MATH70026

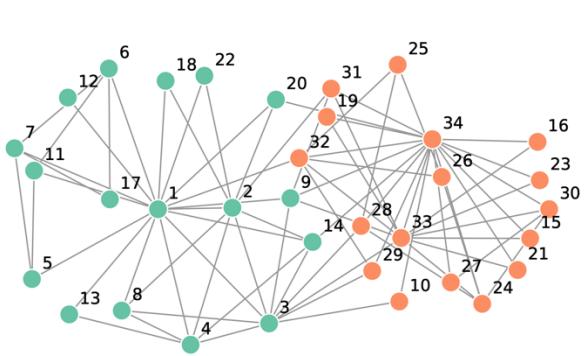
Dr Hardik Rajpal

Any Marvel Fans?

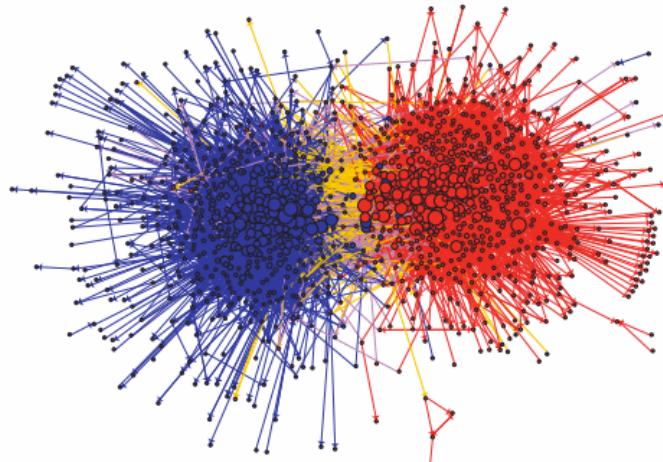
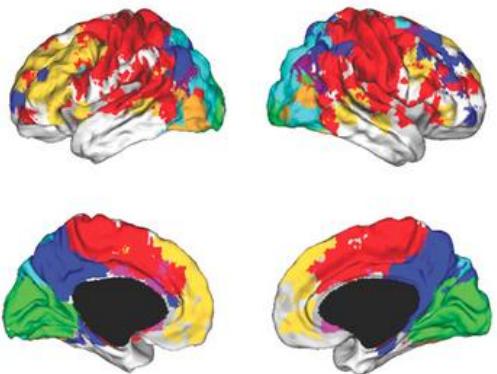
What happened in the Karate Club?



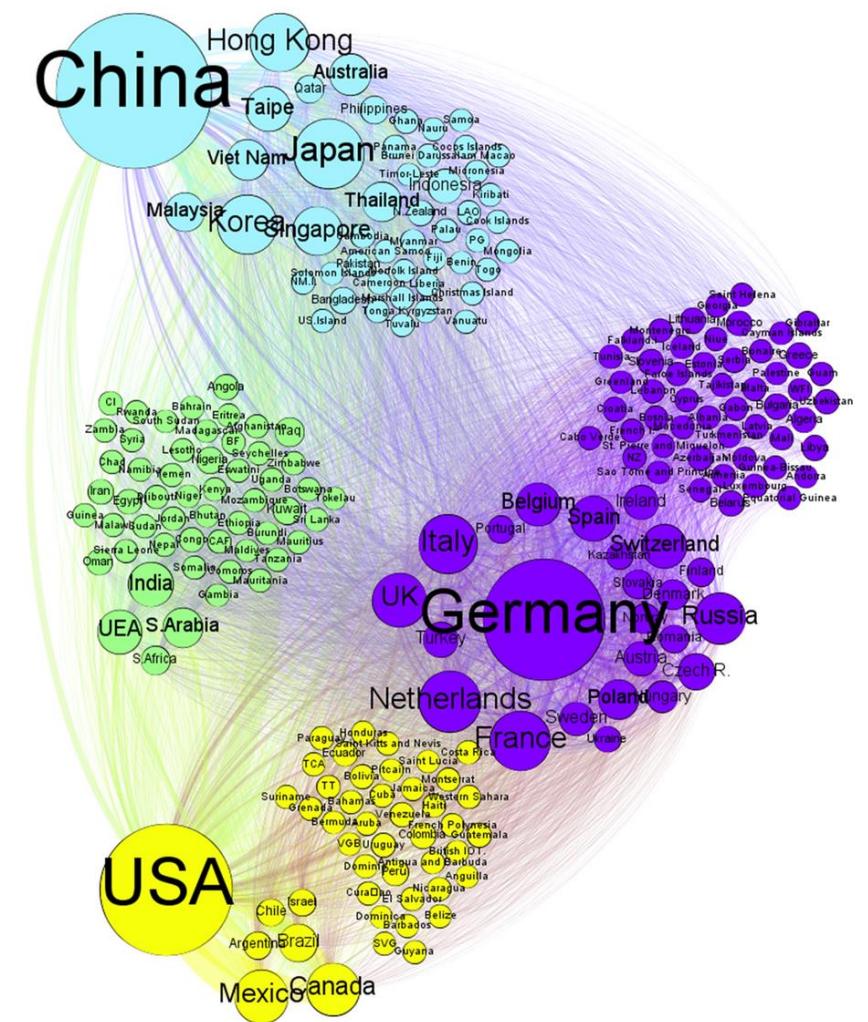
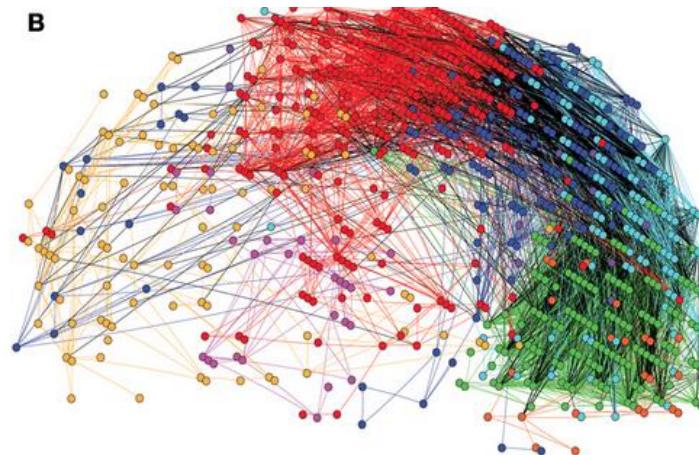
The importance of pairwise relationships



A



B





Join at menti.com | use code 3393 8461

 [Mentimeter](#)

Menti

Graph Based Learning



What other systems can the networks help us understand?



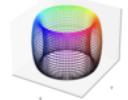
leader bold
creative fast
transpiration focus
inspiration



What other systems can the networks help us understand?
0 responses

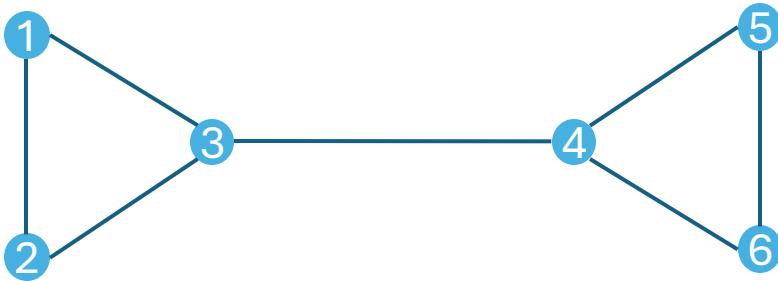
Which method might be better for multiscale data?

0 kNN 0 cSVD 0 Both equally good
0 responses

What is the dimension of the manifold of a torus/doughnut?

0 1D 0 2D 0 3D
0 responses

Who is the most central character in the Marvel Universe?
0 responses

What is a Graph?



Nodes $V = \{i\}_{i=1}^N$

Edges $E = \{(i,j)\}_{i \sim j}$

$$A = \begin{bmatrix} 0 & 1 & 1 & 0 & 0 & 0 \\ 1 & 0 & 1 & 0 & 0 & 0 \\ 1 & 1 & 0 & 1 & 0 & 0 \\ 0 & 0 & 1 & 0 & 1 & 1 \\ 0 & 0 & 0 & 1 & 0 & 1 \\ 0 & 0 & 0 & 1 & 1 & 0 \end{bmatrix} \quad d = \begin{bmatrix} 2 \\ 2 \\ 3 \\ 3 \\ 2 \\ 2 \end{bmatrix}$$

- For undirected graphs , A is symmetric.
- For directed graphs d_{in} and d_{out} are defined separately as,

$$d_{in} = A\mathbf{1} \quad \text{and} \quad d_{out} = A^T\mathbf{1}$$

For weighted graphs,
 $A_{ij} = w_{ij}$, where $w_{ij} \in \mathbb{R} \forall A_{ij} > 0$

What is a Graph Laplacian?

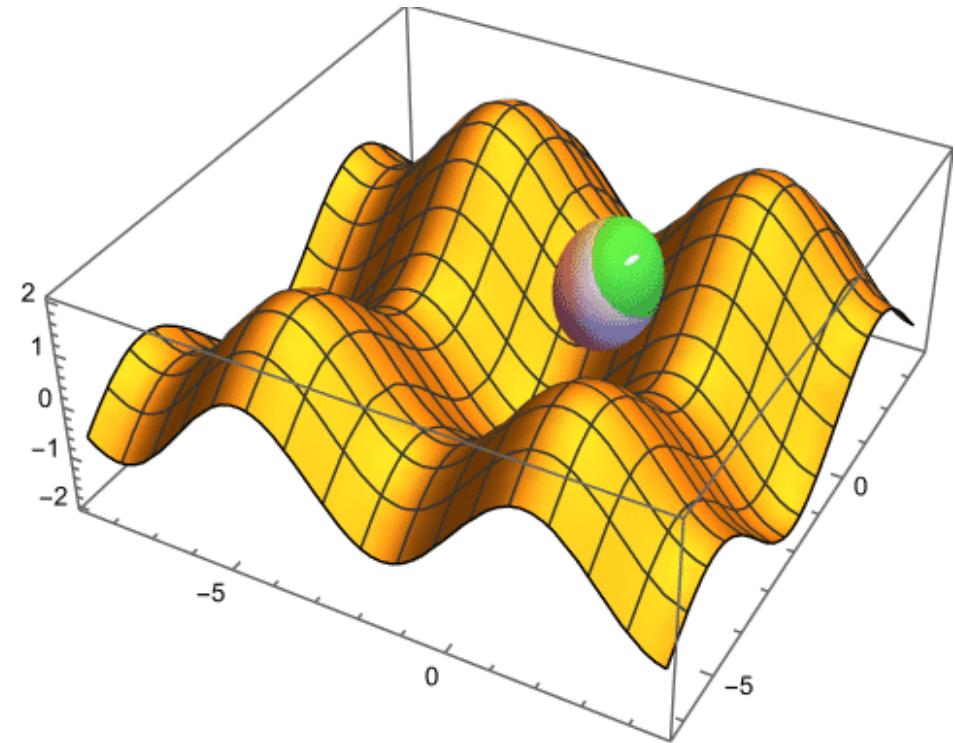
The Laplacian is a second-order differential operator over space, which measures how a quantity changes at a point relative to its surroundings.

In the diffusion equation, it measures how a quantity disperses in space over time.

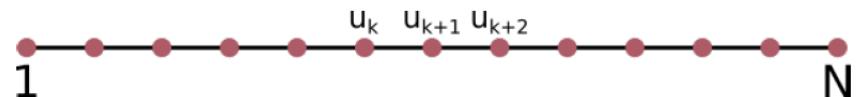
$$u = u(x, t), \quad \frac{\partial u}{\partial t} = \nabla^2 u$$

Laplacian Operator

$$\nabla^2 u = \frac{d}{dx} \left(\frac{d}{dx} u \right)$$



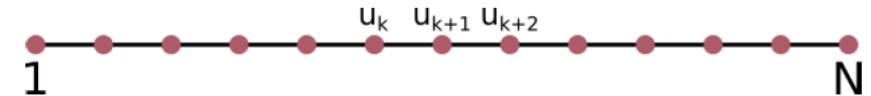
The 1D discrete space can be represented as a graph of nodes along a line



What is a Graph Laplacian?

Now , we can rewrite the diffusion equation in the matrix form on this graph.

$$\frac{d\mathbf{u}}{dt} = \begin{pmatrix} -1 & 1 & & & & 0 \\ 1 & -2 & 1 & & & \\ & \ddots & \ddots & \ddots & & \\ & & 1 & -2 & 1 & \\ & & & \ddots & \ddots & \ddots \\ 0 & & & & 1 & -2 & 1 \end{pmatrix} \begin{pmatrix} u_1 \\ \vdots \\ u_k \\ u_{k+1} \\ u_{k+2} \\ \vdots \\ u_N \end{pmatrix}$$



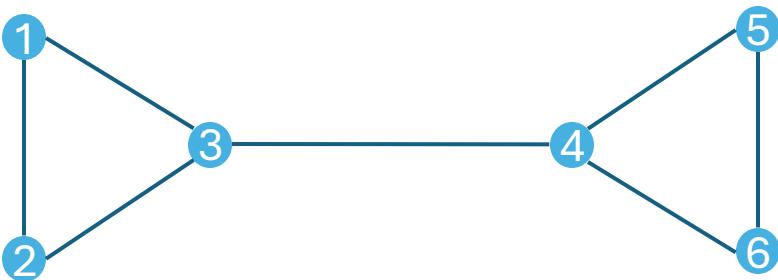
$$A = \begin{pmatrix} 0 & 1 & & & \\ 1 & 0 & \ddots & & \\ & \ddots & \ddots & 1 & \\ & & 1 & 0 \end{pmatrix} \quad d = \begin{pmatrix} 1 \\ 2 \\ \vdots \\ 2 \\ 1 \end{pmatrix}$$

$$\frac{d\mathbf{u}}{dt} = -(diag(d) - A) \mathbf{u} = -\mathcal{L}\mathbf{u}$$

$$\mathcal{L} = diag(d) - A = \mathcal{D} - A$$

\mathcal{L} is known as the combinatorial graph Laplacian

What is a Graph Laplacian?



Nodes $V = \{i\}_{i=1}^N$

Edges $E = \{(i,j)\}_{i \sim j}$

$$A = \begin{bmatrix} 0 & 1 & 1 & 0 & 0 & 0 \\ 1 & 0 & 1 & 0 & 0 & 0 \\ 1 & 1 & 0 & 1 & 0 & 0 \\ 0 & 0 & 1 & 0 & 1 & 1 \\ 0 & 0 & 0 & 1 & 0 & 1 \\ 0 & 0 & 0 & 1 & 1 & 0 \end{bmatrix} \quad d = \begin{bmatrix} 2 \\ 2 \\ 3 \\ 3 \\ 2 \\ 2 \end{bmatrix}$$

$$\mathcal{L} = \text{diag}(d) - A = \mathcal{D} - A =$$

$$\begin{bmatrix} 2 & -1 & -1 & 0 & 0 & 0 \\ -1 & 2 & -1 & 0 & 0 & 0 \\ -1 & -1 & 3 & -1 & 0 & 0 \\ 0 & 0 & -1 & 3 & -1 & -1 \\ 0 & 0 & 0 & -1 & 2 & -1 \\ 0 & 0 & 0 & -1 & -1 & 2 \end{bmatrix}$$

Pairwise relationships in data

We can represent our dataset of N samples as a vector in the p -dimensional space.

In this course, we have explored meaningful metrics that quantify pairwise similarity (or dissimilarity) among data samples.

$$\{\mathbf{x}^{(i)}\}_{i=1}^N \quad \mathbf{x}^{(i)} \in \mathbb{R}^p \quad \mathbf{x}^{(i)} = \begin{pmatrix} x_1^{(i)} \\ \vdots \\ x_p^{(i)} \end{pmatrix}$$

(cosine similarity)

$$S_{ij} = \frac{\mathbf{x}^{(i)} \cdot \mathbf{x}^{(j)}}{\|\mathbf{x}^{(i)}\| \|\mathbf{x}^{(j)}\|}$$

(statistical similarity)

$$S_{ij} = \frac{\text{cov}(\mathbf{x}^{(i)}, \mathbf{x}^{(j)})}{\sqrt{\text{cov}(\mathbf{x}^{(i)}, \mathbf{x}^{(i)})} \sqrt{\text{cov}(\mathbf{x}^{(j)}, \mathbf{x}^{(j)})}}$$

(Euclidean distance)

$$D_{ij} = \|\mathbf{x}^{(i)} - \mathbf{x}^{(j)}\|$$

(kernel matrix)

$$D_{ij} = \exp \left[\frac{-\|\mathbf{x}^{(i)} - \mathbf{x}^{(j)}\|^2}{c} \right]$$

$$X_{N \times p} = \begin{pmatrix} \mathbf{x}^{(1)T} \\ \vdots \\ \mathbf{x}^{(N)T} \end{pmatrix}$$

We can leverage these measures of similarity among data samples to construct a weighted graph of N nodes and edges as the (dis)similarity matrix with adjacency matrices $D_{N \times N}$ or $S_{N \times N}$

Pairwise relationships in data: Quick Aside

In this lecture we will focus on the network of N samples for clustering and dimensionality reduction.

However, in many other applications, the network of descriptors/features is more informative.

Especially when features represent data from different parts of a system.
E.g. Timeseries data from different brain regions.

Statistical interdependence relationships can be quantified amongst the features using Pearson correlation, mutual information, granger causality, etc.

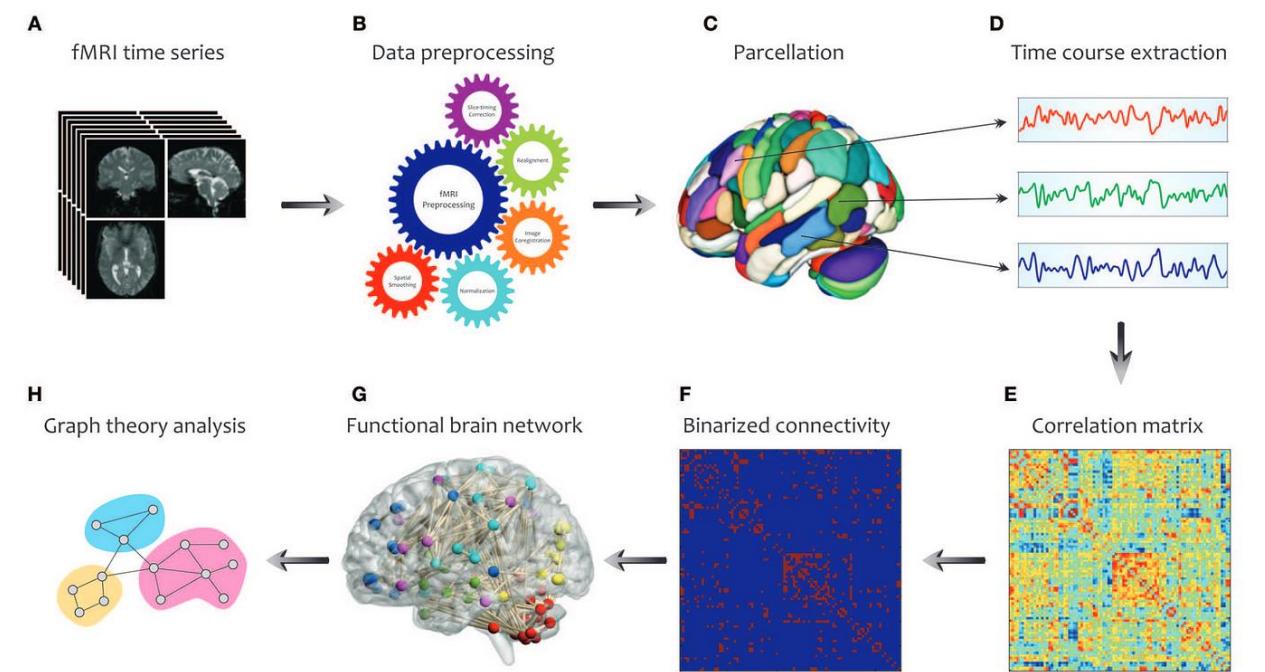
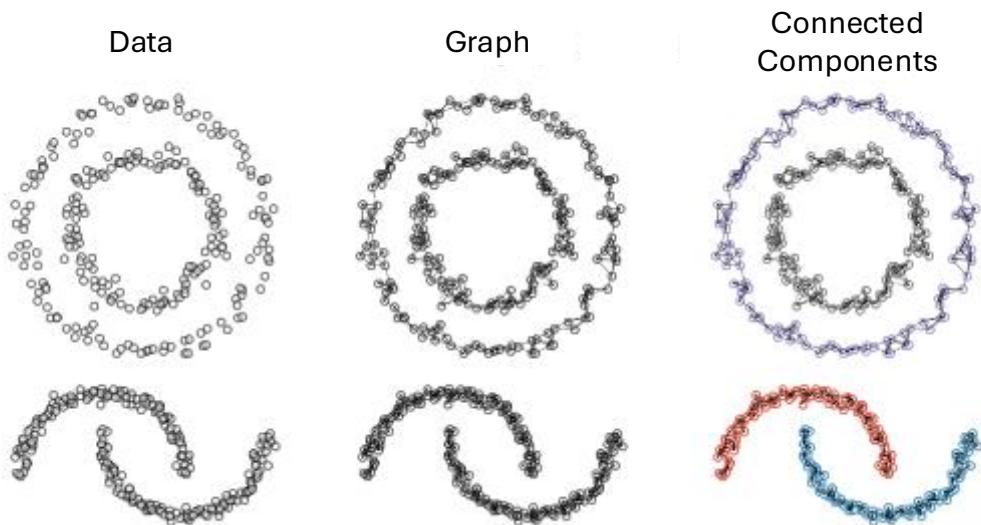


Image from: Touihiri, Mariem. *Decoding Brain Connections: Exploring Neural Function Networks Through fMRI Graph Analysis*, Medium 2023

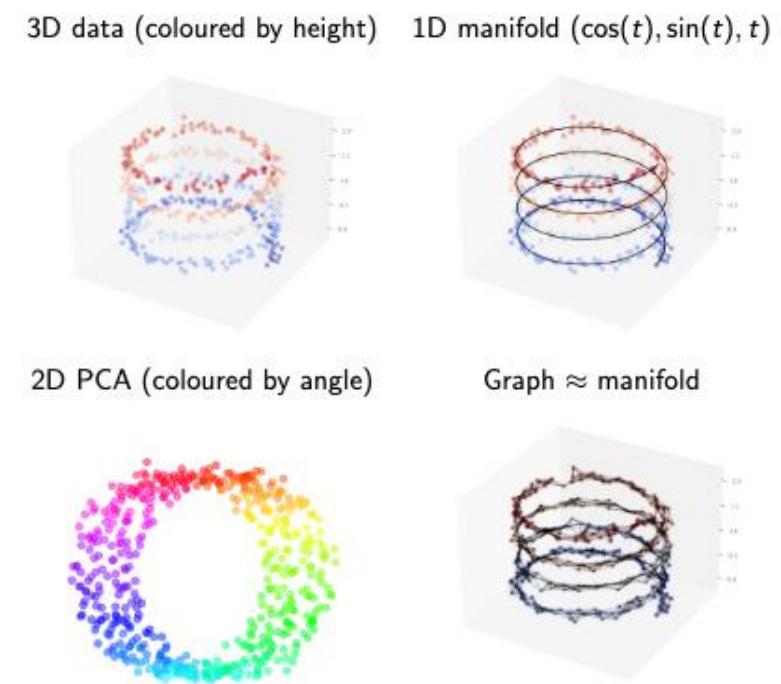
What can graphs tell us?

Clustering



Graphs leverage the intrinsic geometry of the data to facilitate unsupervised learning.

Dimensionality Reduction



Graphs from Data

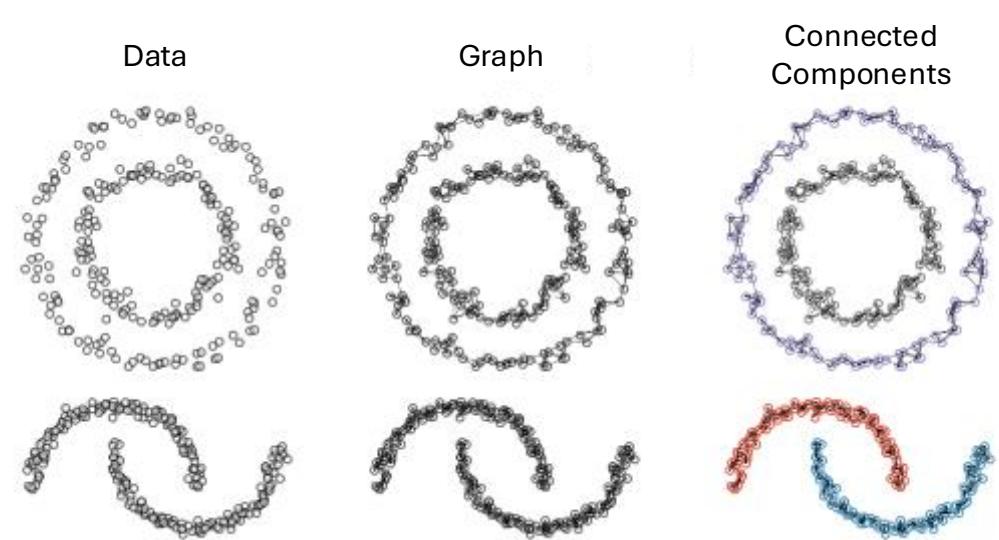
Although the (dis)similarity matrices $D_{N \times N}$ or $S_{N \times N}$, can be used as the adjacency matrix of the graph.

Their fully connected nature hides the underlying geometry, and graph sparsification methods are needed to reduce the density of the graph.

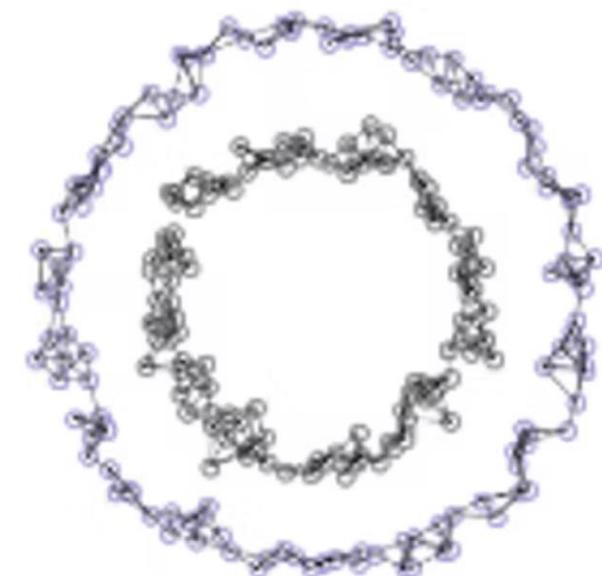
Revealing the underlying geometry.

The most common ones are,

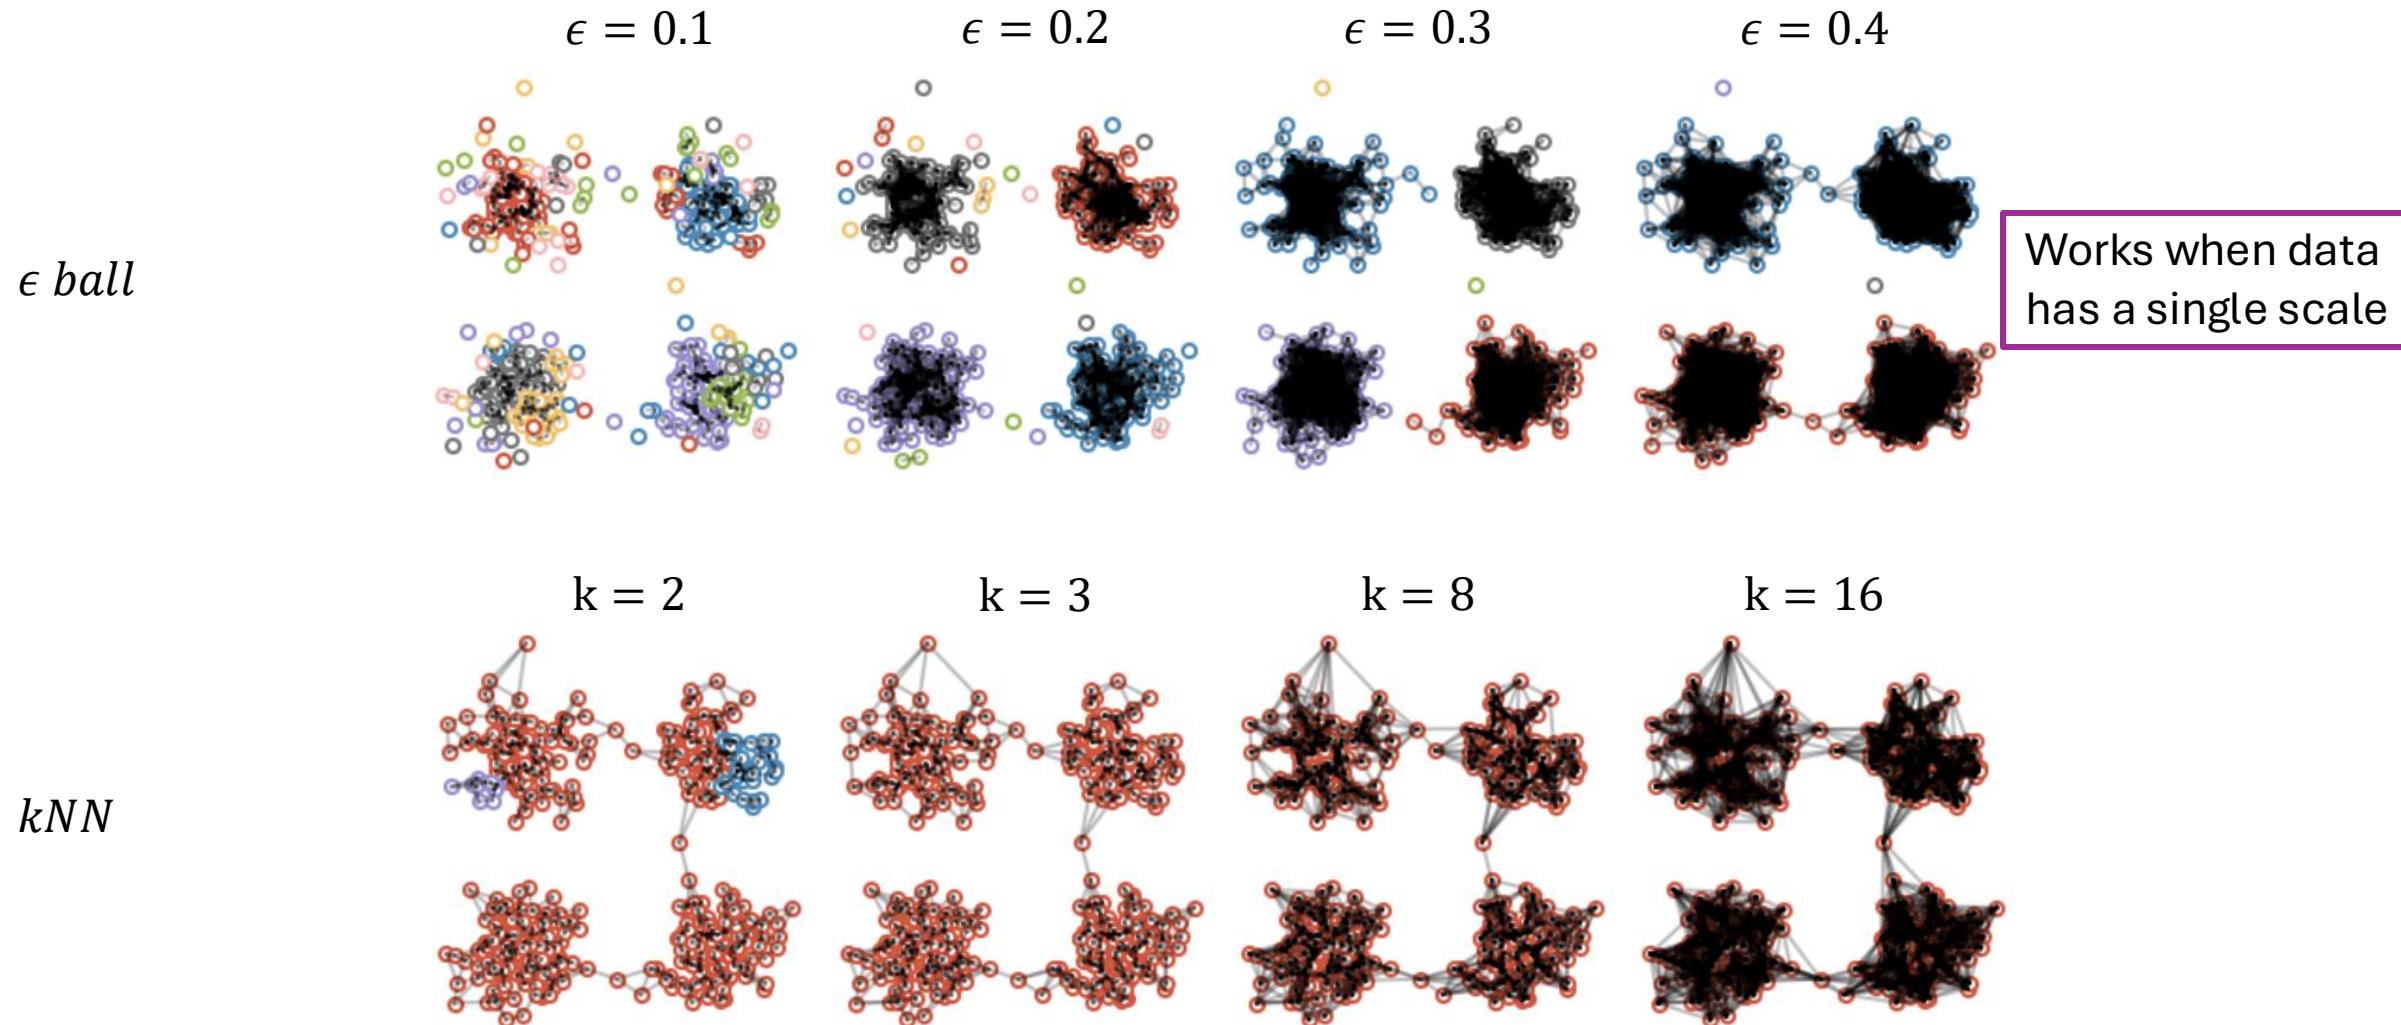
- Thresholding – Removes all edges below a threshold
- ϵ -ball – Nodes within ϵ distance from each other form an edge
- k -NN – The top k nearest neighbours form an edge



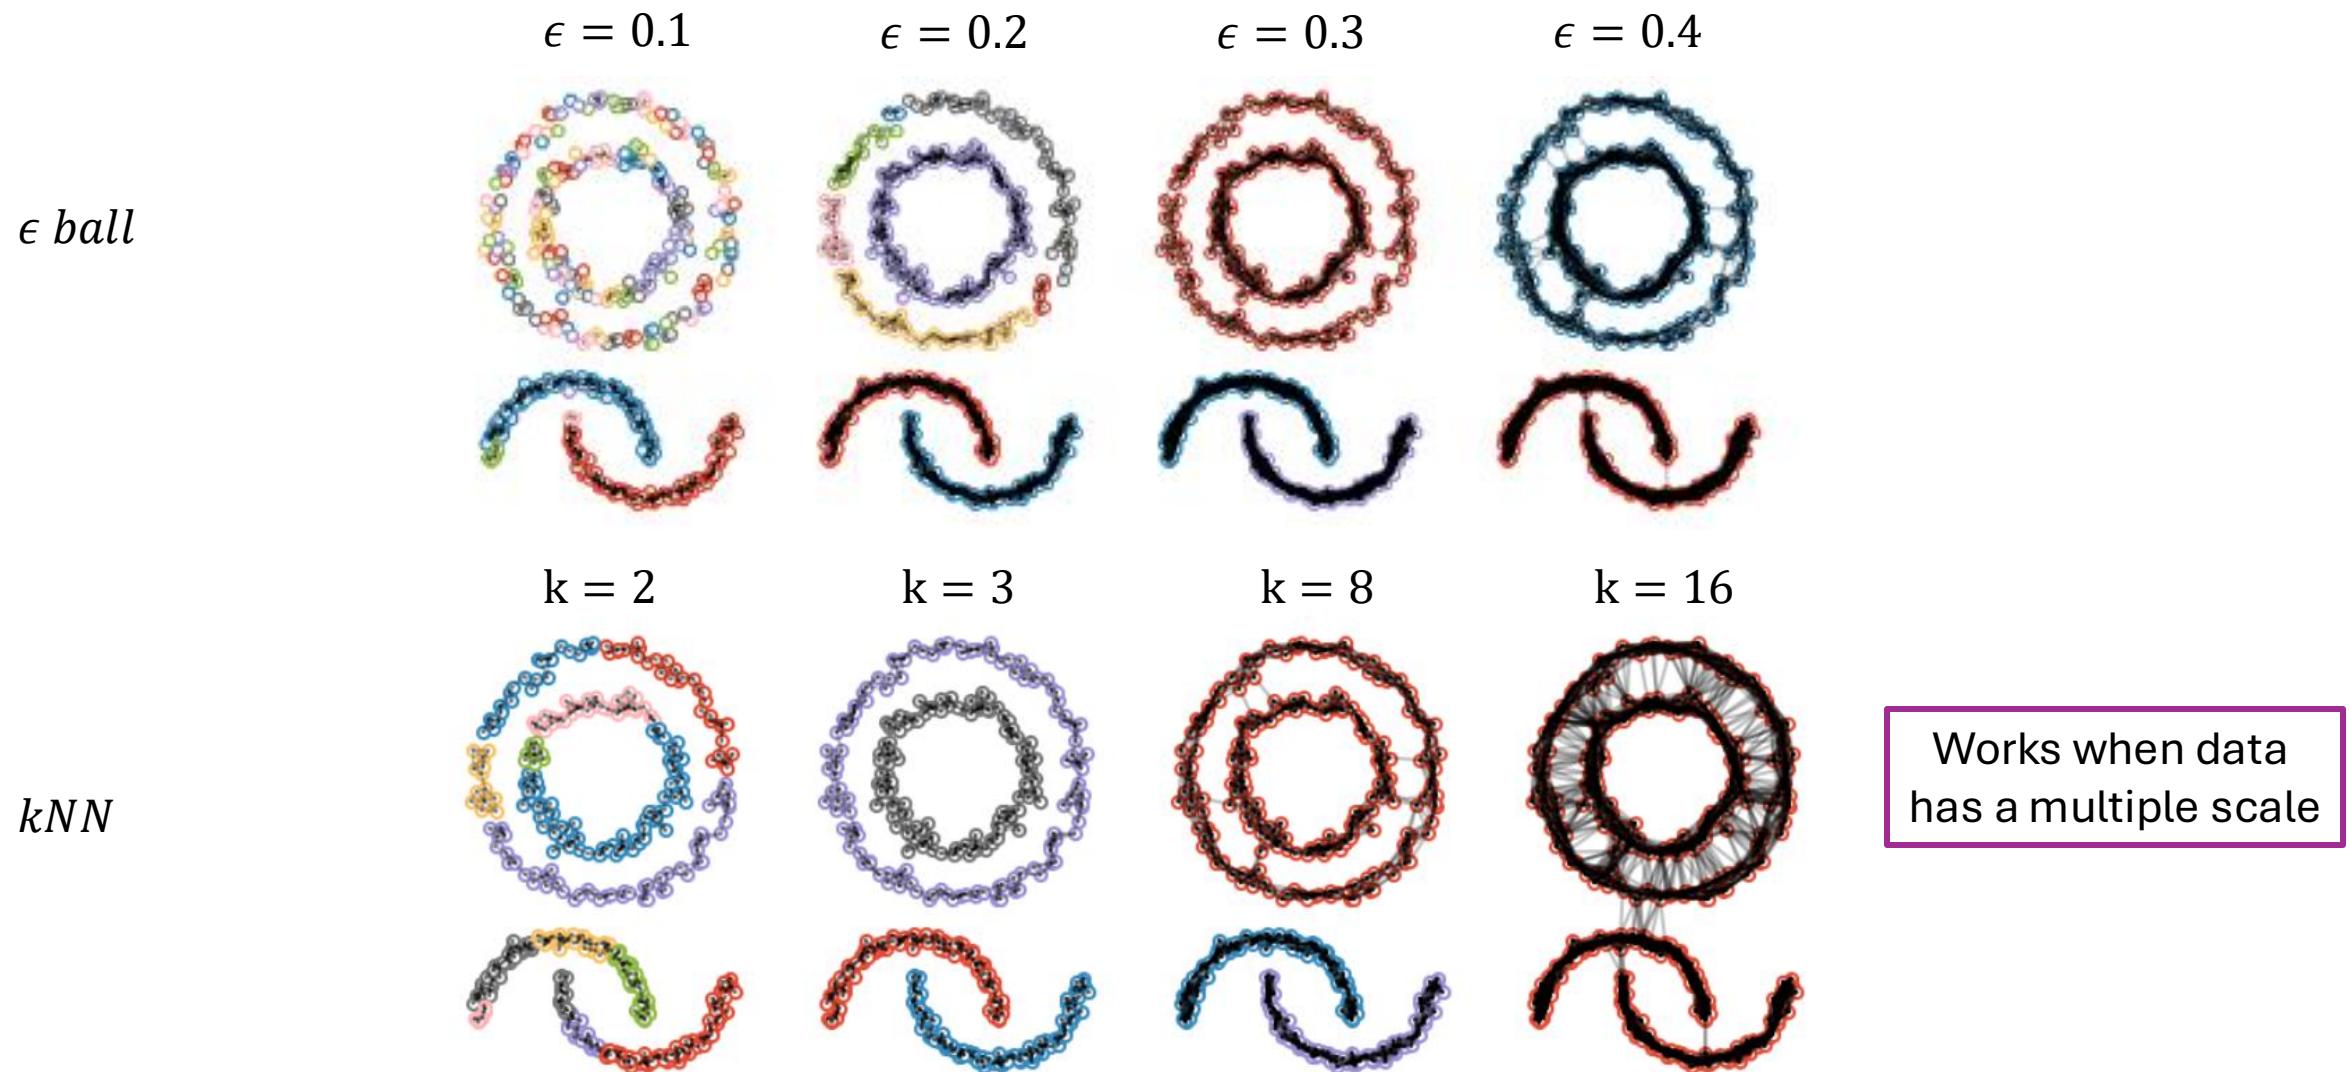
Join at menti.com | use code 3393 8461



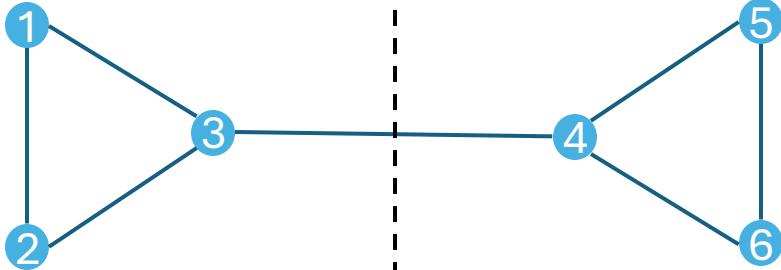
Clustering connected components



Clustering connected components



Clustering on (connected) Graphs



Suppose we wish to bipartition this graph into two clusters, S_1 and S_2 , **minimizing the number of edges between the two clusters**

$$C = \frac{1}{2} \sum_{i \in S_1, j \in S_2} A_{ij}$$

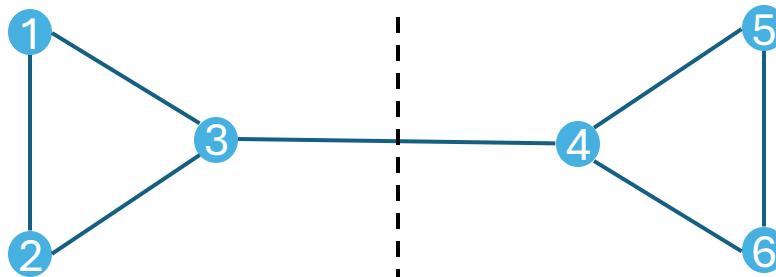
$$s_i = \begin{cases} +1 & \text{if } i \in S_1 \\ -1 & \text{if } i \in S_2 \end{cases} \quad t_{ij} = \frac{1}{2}(1 - s_i s_j) = \begin{cases} 1 & \text{if } i \in S_1, j \in S_2 \\ 0 & \text{otherwise} \end{cases}$$

$$C = \frac{1}{4} \sum_{i,j} t_{ij} A_{ij} = \frac{1}{4} \left(\sum_{i,j} A_{ij} - \sum_{i,j} A_{ij} s_i s_j \right)$$

$$\sum_{i,j} A_{ij} = \sum_i d_i = \sum_i d_i s_i^2 = \sum_{i,j} d_i s_i s_j \delta_{ij}$$

$$C = \frac{1}{4} \sum_{i,j} \left(\underbrace{d_i \delta_{ij}}_{\text{Term 1}} - \underbrace{A_{ij}}_{\text{Term 2}} \right) s_i s_j = \frac{1}{4} \mathbf{s}^T L \mathbf{s}$$

Clustering on (connected) Graphs



The lowest value of C_{min}^* is achieved by the eigenvector v_2 of the graph Laplacian with the second smallest or the first non-trivial eigenvalue λ_2

v_2 is also known as the Fiedler vector.

The number of 0 eigenvalues of the Laplacian gives the number of connected components of the graph.

$$C = \frac{1}{2} \sum_{i \in S_1, j \in S_2} A_{ij} = \frac{1}{4} \mathbf{s}^T L \mathbf{s}$$

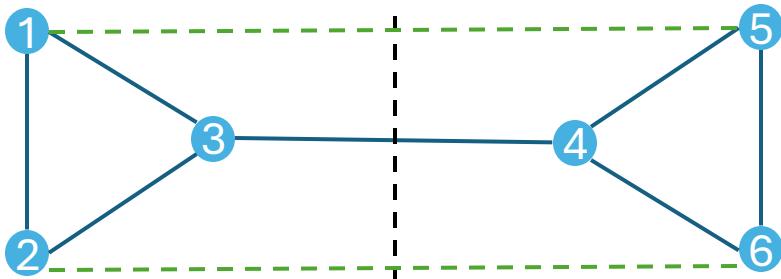
$$\min_{\mathbf{s} \in \{-1, +1\}^N} \mathbf{s}^T L \mathbf{s} \quad \text{under the constraints} \begin{cases} \mathbf{s}^T \cdot \mathbf{s} &= n_1 + n_2 = N \\ \mathbf{s}^T \cdot \mathbf{1} &= n_1 - n_2 \end{cases}$$

$$\mathcal{C}(\mathbf{s}, \lambda, \mu) = \mathbf{s}^T L \mathbf{s} + \lambda(N - \mathbf{s}^T \cdot \mathbf{s}) + 2\mu((n_1 - n_2) - \mathbf{s}^T \cdot \mathbf{1})$$

$$\nabla_{\mathbf{s}} \mathcal{C} \Big|_{\mathbf{s}^*} = 0.$$

$$C_{min}^* = \lambda_2 \frac{n_1 n_2}{N}$$

Clustering on (connected) Graphs



$$\mathcal{L} = \begin{bmatrix} 2 & -1 & -1 & 0 & 0 & 0 \\ -1 & 2 & -1 & 0 & 0 & 0 \\ -1 & -1 & 3 & -1 & 0 & 0 \\ 0 & 0 & -1 & 3 & -1 & -1 \\ 0 & 0 & 0 & -1 & 2 & -1 \\ 0 & 0 & 0 & -1 & -1 & 2 \end{bmatrix}$$
$$\mathcal{L}' = \begin{bmatrix} 3 & -1 & -1 & 0 & -1 & 0 \\ -1 & 3 & -1 & 0 & 0 & -1 \\ -1 & -1 & 3 & -1 & 0 & 0 \\ 0 & 0 & -1 & 3 & -1 & -1 \\ -1 & 0 & 0 & -1 & 3 & -1 \\ 0 & -1 & 0 & -1 & -1 & 3 \end{bmatrix}$$

$$v_2 = (-0.465, -0.465, -0.261, 0.261, 0.465, 0.465)$$

$$v_2 = (0.408, 0.408, 0.408, -0.408, -0.408, -0.408)$$

The sign of each node in the Fiedler vector specifies the partition it belongs to.

$$\lambda_2 = 0.438$$

$$\lambda_2 = 2$$

λ_2 is also known as the *algebraic connectivity* of the graph. It is a measure of how difficult it is to disconnect the graph.

Beyond Bipartitions

- Real networks often have more than 2 clusters.
Here, we discuss two ways of extending the framework using the graph Laplacian.
- ***Sequential Bipartition***
Iteratively identify all partitions of the system
- ***Spectral Clustering*** using more Laplacian Eigenvectors.
 - Consider a $N \times m$ matrix of the m eigenvectors corresponding to the smallest m non-trivial eigenvalues of the graph Laplacian.
 - Each row characterises the N nodes in the m -dimensional eigenspace.
 - Carry out K -means clustering on the row vectors of the $N \times m$ matrix to find K clusters.

Gene Regulatory Network of E. Coli

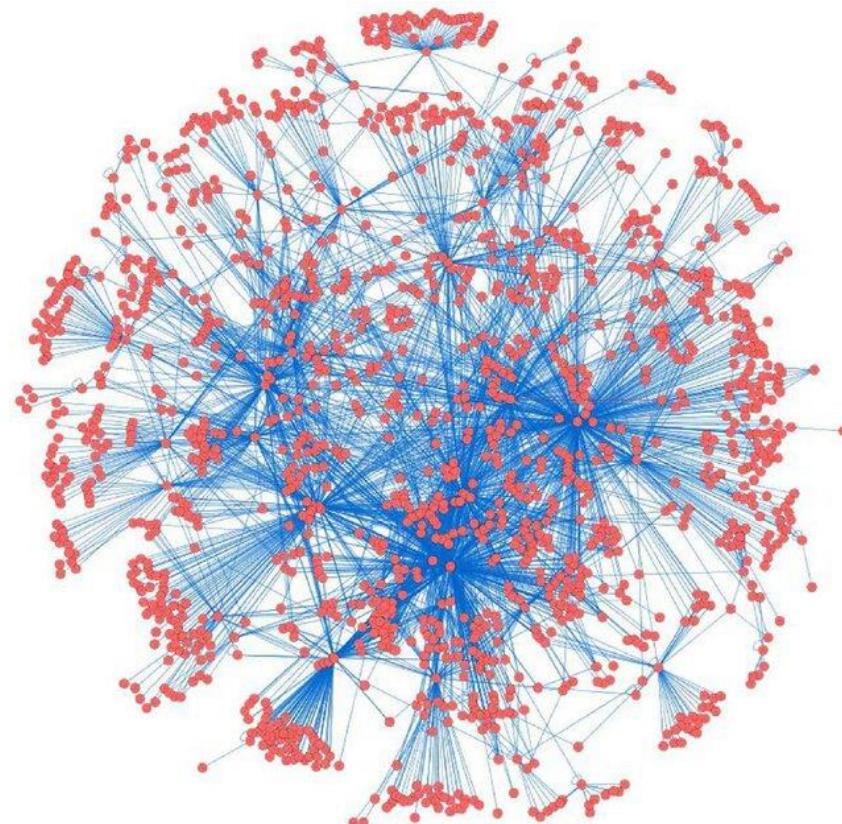
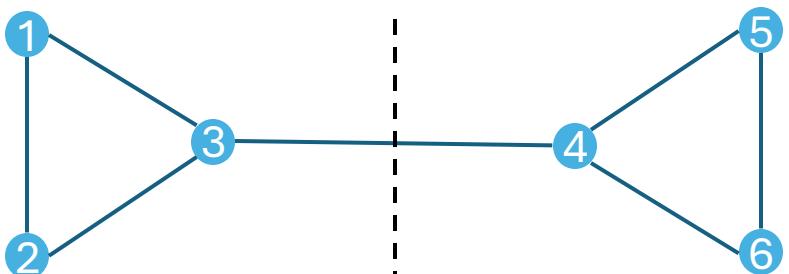


Figure from, Allen, J.D., Xie, Y., Chen, M., Girard, L. and Xiao, G., 2012. Comparing statistical methods for constructing large scale gene networks. *PloS one*

Modularity



$$A = \begin{bmatrix} 0 & 1 & 1 & 0 & 0 & 0 \\ 1 & 0 & 1 & 0 & 0 & 0 \\ 1 & 1 & 0 & 1 & 0 & 0 \\ 0 & 0 & 1 & 0 & 1 & 1 \\ 0 & 0 & 0 & 1 & 0 & 1 \\ 0 & 0 & 0 & 1 & 1 & 0 \end{bmatrix} \quad H = \begin{bmatrix} 1 & 0 \\ 1 & 0 \\ 1 & 0 \\ 0 & 1 \\ 0 & 1 \\ 0 & 1 \end{bmatrix}$$

$$\frac{1}{2} H^T A H = \begin{bmatrix} 3 & \frac{1}{2} \\ \frac{1}{2} & 3 \end{bmatrix}$$

The modularity algorithm aims to cluster the graph by *labelling nodes, making the adjacency matrix **maximally block-diagonal***.

If H is the $N \times k$ membership matrix then,
The total number of edges within the defined clusters
can be written as,

$$\frac{1}{2} \text{Tr}[H^T A H]$$

Along with maximising $\text{Tr}[H^T A H]$ to find the optimal partition, the algorithm also compares this quantity to a null model with the same degree distribution with the adjacency matrix,

$$R = \frac{1}{2E} d d^T$$

Thus, the optimization function for the algorithm can be written as,

$$Q = \frac{1}{2E} \text{Tr} \left[H^T \left[A - \underbrace{\frac{1}{2E} d d^T}_Z \right] H \right]$$

Modularity

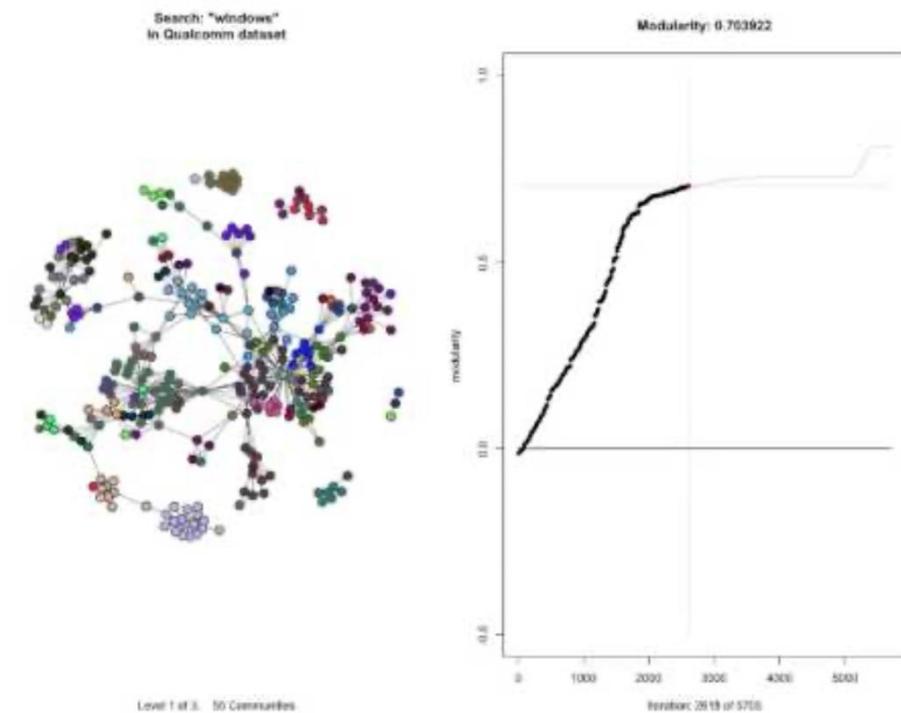
$$Q = \frac{1}{2E} \text{Tr} \left[H^T \underbrace{\left[A - \frac{1}{2E} dd^T \right]}_Z H \right]$$

The Modularity method aims to **maximise** the modularity function Q over all membership matrices H to find the optimal partition

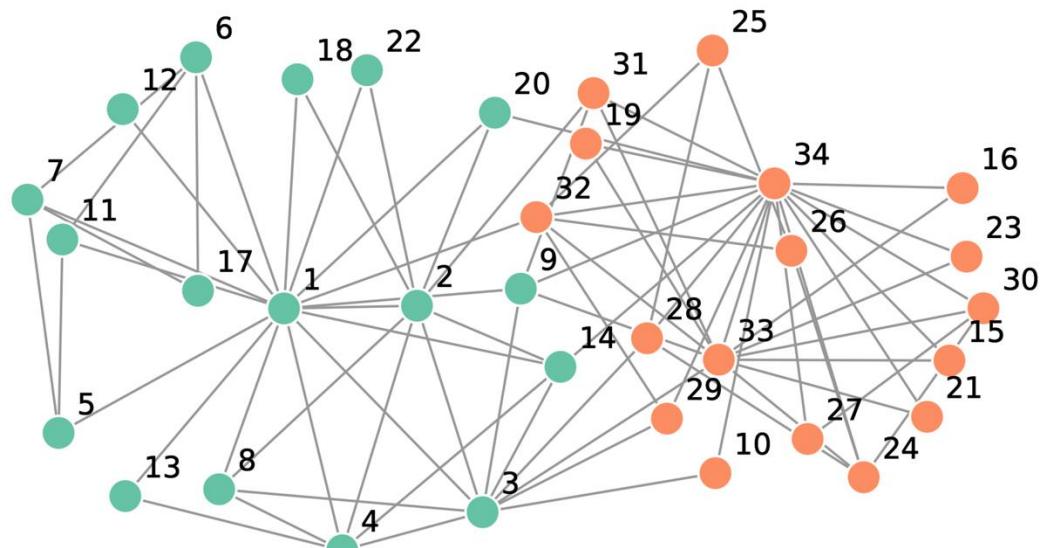
$$\max_H Q$$

Practical algorithms to maximize Q

- Iterative bipartitions of Z using the leading eigenvector until $\Delta Q < 0$
- The *Louvain Algorithm*
A greedy agglomerative algorithm that adds nodes of the graph to a partition until no improvement of Q is achieved.



Spectral and Modularity Bipartitions



```
import networkx as nx

zkc = nx.karate_club_graph()
gt_membership = [zkc.nodes[v]['club'] for v in zkc.nodes()]
```

Fun Exercise

Compare the bipartitions obtained using modularity optimization and using the Fiedler vector v_2 to ground truth partition using the Normalised Variation of Information between any learned partition and the ground truth.

Mutual Information:

$$MI(X, Y) = \sum_{i=1}^K \sum_{j=1}^{K'} P_{XY}(i, j) \log \frac{P_{XY}(i, j)}{P_X(i)P_Y(j)}$$

Clustering entropy

Normalised Variation of Information:

$$NVI(X, Y) = \frac{H(X) + H(Y) - 2MI(X, Y)}{H(X) + H(Y) - MI(X, Y)}$$

$$H(X) = -\sum_{i=1}^K P_X(i) \log P_X(i)$$

and others, see lecture notes

Break?

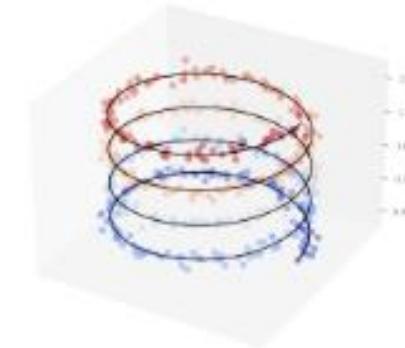
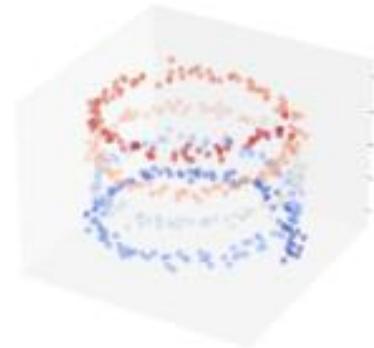
Dimensionality reduction using Graphs

Manifold Assumption: The data lie on some low dimensional manifold of dimensionality m , embedded in a high-dimensional space of dimensionality $p > m$

Manifold is a collection of points such that local neighbourhood of each point resembles a Euclidean space of dimension m

Since graphs capture the interactions between a node and its neighbours connected by edges, **constructing a graph that preserves the local structure of the data can approximate the low dimensional manifold.**

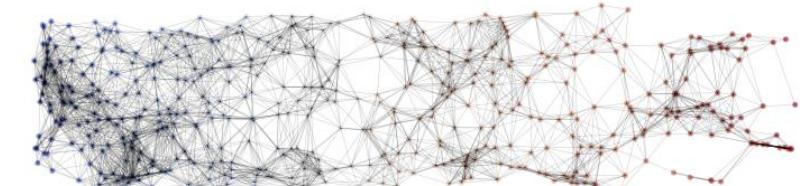
3D data (coloured by height) 1D manifold ($\cos(t), \sin(t), t$)



3D data

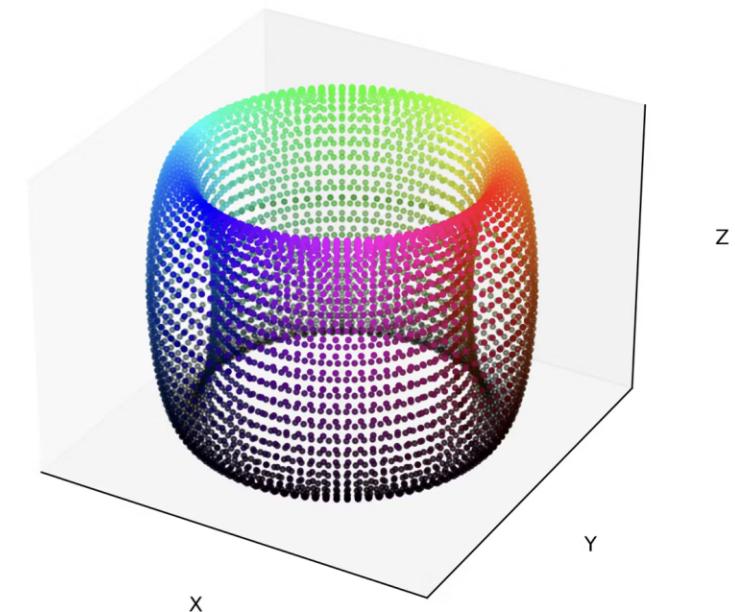


2D Manifold

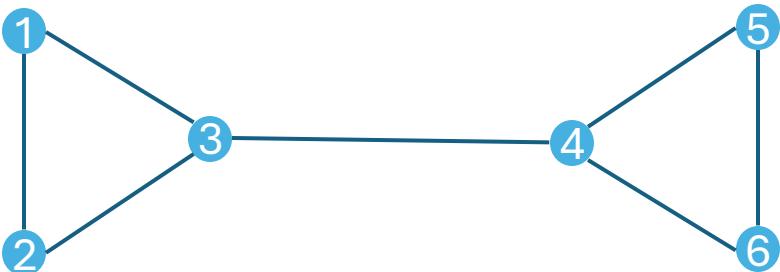


Join at menti.com | use code 3393 8461

 **Mentimeter**



Spectral Embedding



$$v_2 = (-0.465, -0.465, -0.261, 0.261, 0.465, 0.465)$$

While spectral embedding maps neighbouring nodes to a close distance in the manifold, it doesn't ensure nodes far away in the graph are mapped far away in the cost function.

Preserves the local structure of the data.

Which means if i and j are neighbours in the graph then they should be close together in the mapped lower dimensional space.

$$C := \frac{1}{2} \sum_{i,j} \|\hat{\mathbf{x}}^{(i)} - \hat{\mathbf{x}}^{(j)}\|_2^2 A_{ij},$$

$$C = \text{tr}(\hat{X}^T L \hat{X}),$$

Solution: the m eigenvectors of the graph Laplacian where eigenvalues are in ascending order
(ignoring the trivial zero eigenvalues)
First shown in *Belkin and Niyogi (2002)*

Isomap

Preserves the global structure of the data.

Which means distance between i and j in the graph should be similar to the distance between them in the mapped lower dimensional space.

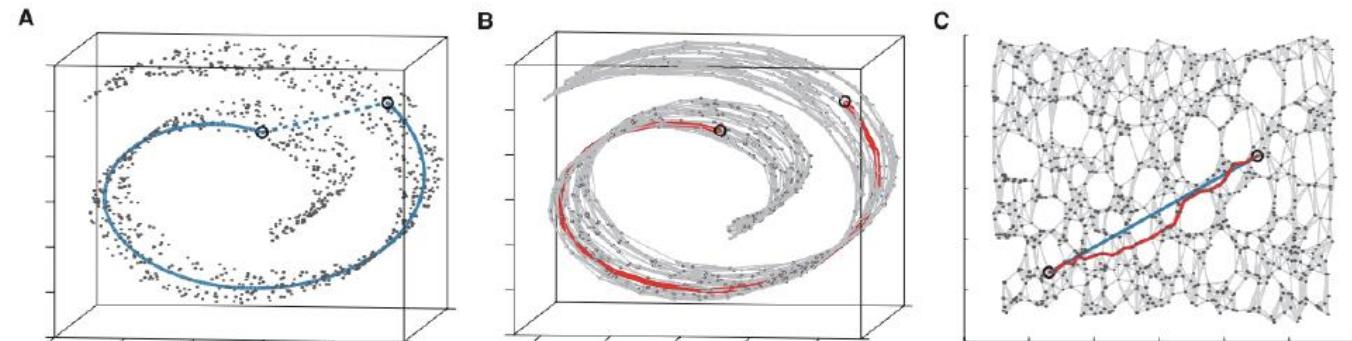
$$C := \sum_{i,j} \left[\|\hat{\mathbf{x}}^{(i)} - \hat{\mathbf{x}}^{(j)}\|_2^2 - d(i,j)^2 \right]^2$$

Where $d(i,j)$ is the geodesic distance or the shortest path length on the graph.

This transformation is equivalent to a kernel PCA where $K_{ij} = -\frac{1}{2}d(i,j)^2$
Tenebaum, de Silva & Langford (2000)

The Isomap algorithm can be summarised as follows,

- Construct a (weighted) graph using k -NN or ϵ -ball, with edge weights as Euclidean distances.
- Compute all pairwise geodesic distances using graph algorithms like Dijkstra's algorithm (for positive edge weights).
- Apply multidimensional scaling (MDS) to the matrix $D_{N \times N}$ of geodesic distances.



Graph Centrality Measures

Graph centrality measures rank nodes based on their level of importance.

Different measures capture different perspectives on what it means to be “important”.

Degree Centrality: Total degree of the node normalised by total number of edges.

$$c_d = \frac{d}{2E}$$

Betweenness Centrality: Total number of shortest paths that pass through a node, normalised by the total number of shortest paths.

Eigenvector Centrality: Nodes that are themselves connected to highly central nodes. The components of the leading eigenvector of the adjacency matrix define the eigenvector centrality of each node.

Page Rank: The algorithm assigns higher centrality to nodes where most random walkers on the graph are likely to end up.

$$\mathbf{c}_{\text{PR } t+1} = \alpha (A\mathcal{D}^{-1}) \mathbf{c}_{\text{PR } t} + (1 - \alpha) \frac{1}{N} \mathbf{1}$$

$$\mathbf{c}_{\text{PR}} = \alpha (A\mathcal{D}^{-1}) \mathbf{c}_{\text{PR}} + (1 - \alpha) \frac{1}{N} \mathbf{1}, \quad \alpha \in [0, 1]$$

Join at menti.com | use code 3393 8461

 Mentimeter

Who is the most central character in the Marvel Universe?

0 responses



Marvel Universe Social Graph

The top ten according to,

Degree

'CAPTAIN AMERICA',
'SPIDER-MAN/PETER PAR',
'IRON MAN/TONY STARK',
'THING/BENJAMIN J. GR',
'MR. FANTASTIC/REED R',
'HUMAN TORCH/JOHNNY S',
'SCARLET WITCH/WANDA',
'WOLVERINE/LOGAN',
'THOR/DR. DONALD BLAK',
'VISION'

Betweenness

'SPIDER-MAN/PETER PAR',
'CAPTAIN AMERICA',
'HAVOK/ALEX SUMMERS',
'IRON MAN/TONY STARK',
'WOLVERINE/LOGAN',
'THING/BENJAMIN J. GR',
'MR. FANTASTIC/REED R',
'THOR/DR. DONALD BLAK',
'HAWK',
'HUMAN TORCH/JOHNNY S',

Eigenvector

'CAPTAIN AMERICA',
'IRON MAN/TONY STARK',
'SCARLET WITCH/WANDA',
'THING/BENJAMIN J. GR',
'SPIDER-MAN/PETER PAR',
'MR. FANTASTIC/REED R',
'VISION',
'HUMAN TORCH/JOHNNY S',
'WOLVERINE/LOGAN',
'BEAST/HENRY & HANK & P'

PageRank

'CAPTAIN AMERICA',
'SPIDER-MAN/PETER PAR',
'IRON MAN/TONY STARK',
'THING/BENJAMIN J. GR',
'MR. FANTASTIC/REED R',
'HUMAN TORCH/JOHNNY S',
'WOLVERINE/LOGAN',
'SCARLET WITCH/WANDA',
'THOR/DR. DONALD BLAK',
'INVISIBLE WOMAN/SUE'



Hero Social Network Data (csv)

6426 Nodes
167207 Edges

Key Takeaways – Graph Based Learning

- Graphs capture pairwise relationships and the underlying geometry of the data to facilitate unsupervised learning.
- Constructing graphs from data using (dis)similarity matrices among samples along with graph sparsification methods like k-NN and ϵ -ball.
- Graph Laplacian plays an important role in both clustering and dimensionality reduction of the data.
- Graphs help uncover the low-dimensional manifold, enabling non-linear dimensionality reduction that preserves the data's local (Spectral Embedding) and global (Isomap) structure.
- Centrality measures provide a way to rank the nodes based on different measures of importances.